

# Towards Verified Faithful Simulation

Vania Joloboff<sup>1</sup>, Jean-François Monin<sup>2</sup>, and Xiaomu Shi<sup>3</sup>

<sup>1</sup> East China Normal University - INRIA - LIAMA

<sup>2</sup> Université de Grenoble - Verimag

<sup>3</sup> Tsinghua University

**Abstract.** This paper presents an approach to construct a verified virtual prototyping framework of embedded software. The machine code executed on a simulated target architecture can be proven to provide the same results as the real hardware, and the proof is verified with a theorem prover. The method consists in proving each instruction of the instruction set independently, by proving that the execution of the C code simulating an instruction yields an identical result to that obtained by a formal executable model of the processor architecture. This formal model itself is obtained through an automated translation process from the architecture specifications. Each independent proof draws a number of lemmas from a generic lemma library and also uses the automation of inversion tactics in the theorem prover. The paper presents the proof of the ARM architecture version 6 Instruction Set Simulator of the Sim-SoC open source simulator, with all of the proofs being verified by the Coq proof assistant, using automated tactics to reduce manual proof development.

## 1 Introduction

In many embedded systems applications nowadays, virtual prototyping is used to design, develop and test new applications. Most of these virtual prototypes include an Instruction Set Simulator (ISS) to simulate the target processor. The ISS runs the target executable binary code in emulating the hardware and generate the outputs that the executable should produce when run on the target platform. An ISS can be used for example to optimize algorithms such as cryptographic software, or to debug new compiler developments, or in the design of many embedded systems applications. Instead of using real hardware prototypes, simulated platforms are more convenient and less expensive. Then, it is important to be sure that the simulator used is faithful to the hardware that it emulates. A *faithful ISS* must produce exactly the same results as the executable would if run on hardware implementation of the instruction set specification, and this guarantee must be proven.

The purpose of our work is to formally verify that the execution of a program on our Instruction Set Simulator for the target ARM architecture indeed produces the expected results, to be certain that the data output from the simulator, the final processor and memory states are indeed identical to the result

obtained with the real hardware. This requires sequential steps, to prove first that the translation from the C code of the simulator to the simulation machine is correct, and second that the simulation of the target machine code is also correct, that is, it preserves the semantics of the computer architecture, together with the fact that all of these proofs are verified using a theorem prover, or proof checker, not subject to human error in the proof elaboration or verification.

The next sections of the paper are organized as follows. Section 2 reviews related work. Section 3 describes the tools that we have used, in particular the CompCert C compiler, a certified compiler for the C language, the Coq proof assistant, and the SimSoC simulator in which our work is integrated. Section 4 presents our contribution to prove the correctness of an ARM Instruction Set Simulator, integrated within SimSoC. In summary, the method consists in proving each instruction of the instruction set independently, by proving that the execution of the C code simulating an instruction yields identical result to that obtained by a formal executable model of the architecture. Each independent proof requires using a number of lemmas from a generic lemmas library and usage of a new inversion tactics in the theorem prover. Finally, our conclusion mentions lessons learned and directions for future work.

## 2 Related Work

Program certification has to be based on a formal model of the program under study. Such a formal model is itself derived from a formal semantics of the programming language. Axiomatic semantics and Hoare logic have been widely used for proving the correctness of programs. For imperative programming languages such as C, a possible approach is to consider tools based on axiomatic semantics, like **Frama-C** [?], a framework for a set of interoperable program analyzers for C. Most of the modules integrated inside rely on ACSL (ANSI/ISO C Specification Language), a specification language based on an axiomatic semantics for C.

**Frama-C** software leverages off from **Why** technology [?,?], a platform for deductive program verification, which is an implementation of Dijkstra’s calculus of weakest preconditions. **Why** compiles annotated C code into an intermediate language. The result is given as input to the VC (Verification Conditions) generator, which produces formulas to be sent to both automatic provers or interactive provers like Coq.

In our case of verifying an instruction set implementation, we have to deal with a very large specification including complex features of the C language. A framework is required that is rich enough to make the specification manageable, using abstraction mechanisms for instance, and in which an accurate definition of C features is available. As we need to verify specific properties referring to a formal version of the ARM architecture, operational semantics offer a more concrete approach to program semantics as it is based on states. The behavior of a piece of program corresponds to a transition between abstract states. This transition relation makes it possible to define the execution of programs by a mathematical computation *relation*. This approach is quite convenient for prov-

ing the correctness of compilers, using operational semantics for the source and target languages (and, possibly intermediate languages). Operational semantics are used in **CompCert** (described below) to define the execution of C programs, or more precisely programs in the subset of C considered by the **CompCert** project. The work presented in this paper is based on this approach. Interesting examples are given by Brian Campbell in the CerCo project [?], in order to show that the evaluation order constraints in C are lax and not uniform.

A very significant verification work has been done to prove the SEL4 operating system[?]. It is comparable to our work in that they have considered a C implementation. The main difference is that they have not considered operational semantics of C, but deduced the proof obligations from the C code, considering the compiler and the architecture as correct. In our work, we believe that the subset of C accepted by **CompCert** is even larger than the subset accepted in SEL4.

Regarding formalization and proofs related to an instruction set, a Java byte code verifier has been proved by Cornelia Pusch[?], the Power architecture semantics has been formally specified in [?], and closer to our work, the computer science laboratory in Cambridge University has used HOL4 to formalize the instruction set architecture of ARM [?]. The objective of their work was to verify an implementation of the ARM architecture with *logical gates*, whereas we consider a ARM architecture simulator coded in C. Reusing the work done at Cambridge in [?] was considered. But, because we need a certified C compiler and our approach is based on **CompCert** C, which is itself coded in Coq, it would have required us to translate all of the C operational semantics as well, which would have been error prone, not to mention the very large effort. It was more convenient to develop our formal model and our proofs in Coq.

Our work is based on the SimSoC simulation framework [?], available as open source software at <http://gforge.inria.fr/projects/simsoc>, described in the next section.

## 3 Background

### 3.1 Coq

Coq [?] is an interactive theorem prover, implemented in OCaml. It allows the expression of mathematical assertions, mechanically checks proofs of these assertions, helps to discover formal proofs, and may extract a certified program from the constructive proof of its formal specification. Coq can also be presented as a dependently typed  $\lambda$ -calculus (or functional language). For a detailed presentation, the reader can consult [?] or [?]. Coq proofs are typed functions and checking the correctness of a proof boils down to type-checking.

The logic supported by Coq includes arithmetic, therefore it is too rich to be decidable. As full automation is not possible for generating proofs, human interaction is essential. The latter is realized by *proof scripts*, which are sequences of commands for building a proof step by step. Coq also provides built-in *tactics*

implementing various decision procedures for suitable fragments of the calculus of inductive constructions and a language which can be used for automating the search of proofs and shortening scripts.

When a proof has been interactively developed, Coq automatically verifies the proof, or possibly signals where errors are located. Our work has consisted in developing proofs demonstrating that the C functions simulating the behavior of the ARM processor indeed implement the ARM architecture semantics.

### 3.2 Compert-C

**CompCert** is a formally verified compiler for the C programming language provided by INRIA [?,?], which currently targets Power, ARM and 32-bit x86 architectures. The compiler is specified, programmed, and proved in Coq. It aims to be used for programming embedded systems requiring high reliability. The generated assembly code is proved to behave exactly the same as the input C program, according to a formally defined operational semantics of the language.

A key point is that we are considering here C programs compliant with the definition of ISO-C 99 standard of *correct C programs*. Indeed the ISO-C standard identifies many constructions that are syntactically correct, but have undefined semantics such as `a[i++] = i;`. The document identifies about one hundred such constructions, and says that a C compiler in that case basically may choose its own interpretation of the abstract syntax, resulting in *unspecified behavior*. This is very important in our work. All of the C code implementing the ISS is correct with respect to the ISO C standard, meaning that it does not contain any construction with unspecified behavior. **CompCert-C** does not accept such ill-defined expressions and only well formed programs can be translated according to the formal, unambiguous, semantics. All of the C code considered here has unique and well defined semantics. We need to prove that it implements the ARM semantics, but we do not need to worry about multiple interpretations.

Three parts of **CompCert C** are used in this work. The first is that we use the correct machine code generated by the C compiler. The second is the C language operational semantics in Coq from which we get a formal model of the program. Third, we use the **CompCert** Coq library for words, half-words, bytes etc., and bitwise operations to describe the instruction set model. These low level functions have been proven already in **CompCert**, so we can safely re-use them.

It must be noted that the C code of an ISS does not use functions from the C library that invoke the operating system, such as `gettimeofday()`. It uses a very limited number of functions from the C library such as `memset()` or `memcpy()`. **CompCert** provides the formalized properties of such built-in external functions, so we can reason formally on their potential side effects in our proofs.

### 3.3 SimSoC

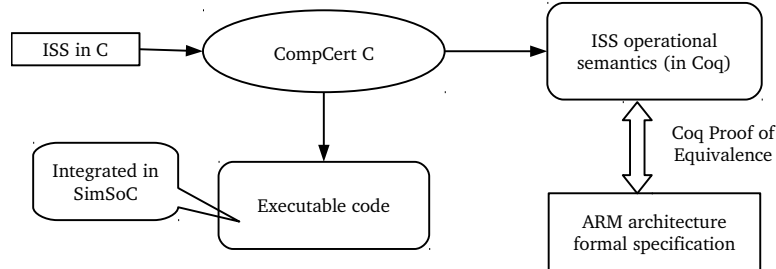
There is abundant literature covering Instruction Set Simulation. Using *interpretive simulation*, such as used in Insulin [?], each instruction of the target

program is fetched from memory, decoded, and executed. With *static translation*, the target application program is decoded at compile time and translated into a new program for the simulation host. The simulation speed is vastly improved [?], but it is not suitable for application programs that generate, or dynamically load code at run-time. Most ISS'es today use some kind of *dynamic binary translation*, initiated with systems such as Embra [?].

As mentioned above, the target ISS for the verification is integrated within **SimSoC** [?], a full system simulator of System-on-Chips, available as open source software. **SimSoC** takes as input real binary code and executes simulation models of the complete embedded system: processor, memory units, interconnect, and peripherals. The chip simulator also includes a network controller simulator, so that the simulator can communicate with the real world. Our proof assumes the existence of a correct decoder to dynamically generate the translation of the input binary into C structures, e.g. the program that takes the binary input sequence and translates it into a sequence of qualified instructions. It is out of scope of the proof.

**SimSoC** uses the SystemC kernel to simulate hardware parallelism and transaction level modeling (TLM) to model communications between the modules. It includes ISS'es to execute embedded applications on various processors. We are considering here the ARM Version 6 ISS. **SimSoC** supports two modes of dynamic translation. In the first mode, our verification target, the binary decoder translates each instruction into a C structure that has a *semantics function* [?]. It is these C semantic functions that we are verifying here.

## 4 Verified Simulation



**Fig. 1.** Overall goal

The general objective is to obtain a verified simulator is illustrated in Figure 1. Considering the ARM architecture, we need to have the following:

- a formal model of the ARM instruction set.
- an instruction set simulator of the ARM arcchitecture coded in the (**CompCert**) C programming language.

- a formal operational semantics of the ISS. As shown in Figure 1, from the ISS source code in C, we can obtain through `CompCert` C on one hand the Coq formal semantics of the compiled C program constructed by `CompCert`, since the intermediate representation of the C compiler is a Coq representation and, on the other hand, the verified machine code, which conforms to this operational semantics as guaranteed by `CompCert`. We use both, the compiled machine code to run simulations, and the formal semantics for the proof.
- prove, using the Coq proof assistant, that the resulting ISS semantics indeed implement the formal model of the ARM processor, which boils down to verifying that the semantics of the simulator accurately modifies the processor (and memory) state representation at each step and ends up in results that comply with the formal model of the ARM architecture.

These steps are described in the following paragraphs.

#### 4.1 Constructing the formal model

Ideally the formal specification of the ARM architecture should be provided by the vendor. But it is not the case, an issue already raised in the work with HOL4 mentioned above [?]. We decided to derive the formal model of ARM architecture in Coq from the architecture reference manual as output of a semi-automated process. The main relevant chapters of the manual are:


- **Programmer’s Model** introduces the main features in ARMv6 architecture, the data types, registers, exceptions, etc;
- **The ARM Instruction Set** explains the instruction encoding in general and puts the instructions in categories;
- **ARM Instructions** lists all the ARM instructions in ARMv6 architecture in alphabetical order and the **ARM Addressing modes** section explains the five kinds of addressing modes.

There are 147 ARM instructions in the ARM V6 architecture. For each instruction, the manual provides its encoding table, its syntax, a piece of pseudo-code explaining its own operation, its exceptions, usage, and notes. Three kinds of information are extracted for each ARM operation: its binary encoding format, the corresponding assembly syntax, and the instruction semantics, which is an algorithm operating on the processor state. This algorithm may call basic functions defined elsewhere in the manual, for which we provide a Coq library defining their semantics. Other than these extracted data files, there is still useful information left in the document which cannot be automatically extracted, such as validity constraints information required by the decoder generator. However, the most tedious (then, arguably, error prone) part is described using fairly simple, precise and regular pseudo-code, allowing us to extract the Coq formal model in three automated steps: (i) extracting information from the `.pdf` file; (ii) parsing the data into abstract syntax trees (iii) automated translation from the abstract syntax into Coq formal model.

During this process, a dozen documentation problems were found but none that were relevant to instruction semantics. These documentation mistakes have been acknowledged by ARM Ltd. Moreover, a single mistake in our automated extractor would impact the formal model of many or even all instructions and then become rather easy to detect. The model has then tested on real programs to verify that we obtain the same results, which gives reasonable confidence in the model.

## 4.2 Proof Structure

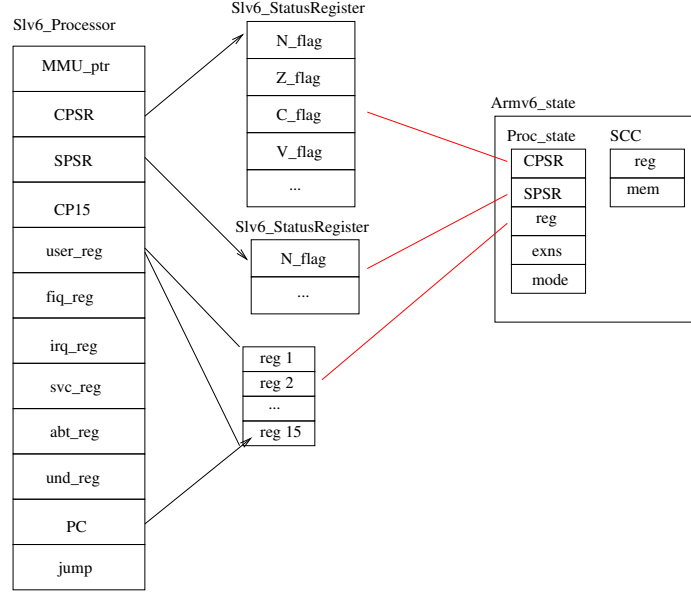
The proof starts from an ISS coded in C, where each instruction is coded as a C function that modifies the processor state and possibly the memory state (but everything is represented in memory on the simulation host machine). Each C function may also call basic functions from a library. As mentioned above, this C code does not include any construction with “unspecified behavior” of the C language specification. To prove that the simulator is correct, we need to prove that, given the initial state of the system, the execution of an instruction as implemented by a C function results in the same state as the formal specification. To establish the proof, a formal model of that C implementation is provided by **CompCert**, which defines operational semantics of C formalized in Coq.



**Fig. 2.** Theorem statement for a given ARM instruction

The proof shall demonstrate that the operational semantics of the C code corresponds to the ARM formal specification. The complete proof is too lengthy for this article, and we only provide here an outline of the method. The state of the ARM V6 processor defined in the formal model is called the *abstract state*. Alternatively, the same state is represented by the data structures corresponding to C semantics that we shall call the *concrete state*. In order to establish correctness theorems we need to relate these two models. Executing the same instruction on the two sides produces a pair of new processor states which should be related as equivalent. Informally, executing the same instruction on a pair of equivalent states should produce a new pair of equivalent states, as schematized by Figure 2. Equivalent states are defined according to a suitable projection from the C concrete state to the abstract model, as represented in Figure 3.

This projection constructs a formal structure from the concrete one. The formal structure obtained should be identical to that obtained through the formal model, otherwise the C code is incorrect.



**Fig. 3.** Projection

### 4.3 Projection

In order to achieve a high speed simulation, the C ISS includes optimizations. In particular, processor state representation in the C implementation is complex, not only due to the inherent complexity of the C language memory model, but also because of optimization and design decisions targeting efficiency. Despite the complexity of the C memory model, the **CompCert** C semantics makes it possible to define and prove the projection. Fortunately, all of the instructions operate on the processor state and there is a single representation of that state in the simulator. It is necessary and sufficient to prove the projection for each variant case of the representation structure. For example, the projection of a register performs a case analysis on possible values, whereas the projection of saved data upon exceptions depends on the type of exception modes. Although there are a number of specific cases to handle, the proof of the projection is relatively straightforward. In more detail:

- The C implementation uses large embedded *structs* to express the ARM processor state. Consequently the model of the state is a complex Coq record



- type, including not only data fields but also proofs to verify access permission, next block pointer, etc.
- Transitions are defined with a relational style (as opposed to a functional style where reasoning steps can be replaced by computations). Relational style is more flexible, especially when dealing with constraints; and fits well with operational semantics.
- The global state is based on a memory model with load and store functions that are used for read/write operations.

The proofs for instructions start from the abstract state described by the formal specification. To verify the projection of the original state, we need the following data: the initial memory state, the local environment, and the formal initial processor state. The projection is meaningful only after the C memory state is prepared for evaluating the current function body representing a ARM instruction. In the abstract Coq model, we directly use the processor state `st`. But on the C side, the memory state is described by the contents of several parameters, including the memory representation of the processor state. We also need to observe the modifications of certain memory blocks corresponding to local variables.

The semantics of **CompCert** C considers two environments. The global environment *genv* maps global function identifiers, global variables identifiers to their blocks in memory, and function pointers to a function definition body. The local environment *env* maps local variables of a function to their memory blocks reference. It maps each variable identifier to its location and its type, and its value is stored in the associated memory block. The value associated to a C variable or a parameter of a C function is obtained by applying `load` to the suitable reference block in memory. These two operations are performed when a function is called, building a local environment and an initialized memory state. When the program starts its execution, *genv* is built. The local environment *env* is built when the associated function starts to allocate its variables. Therefore, on the concrete side, a memory state and a local environment is prepared initially using two steps. First, from an empty local environment, all function parameters and local variables are allocated, resulting into some memory state and the local environment. Second, function parameters are set up using a dedicated function `bind_parameters` and the initial state is thus created.

#### 4.4 Lemmas Library

Next, we need to consider the execution of the instruction. In the C ISS, there is a standalone C function for each ARM V6 instruction. Each function (instruction) has its own correctness proof. Each function is composed of its return type, arguments variables, local variables, and the function body. The function body is a sequence of statements including assignments and expressions. Let us consider as an example the ARM instruction BL (**Branch and Link**). The C code is:

```
void B(struct SLv6_Processor *proc,
      const bool L,
```

```

    const SLv6_Condition cond,
    const uint32_t signed_immed_24){
if (ConditionPassed(&proc->cpsr, cond)){
  if ((L == 1))
    set_reg(proc, 14, address_of_next_instruction(proc));
    set_pc_raw(proc, reg(proc, 15) + (SignExtend_30(signed_immed_24) << 2));
  }
}

```

**CompCert** has designed semantics for **CompCert C** in big-step inductive types for evaluating expressions, which we reuse for the proof. The semantics is defined as a relation between an initial expression and an output expression after evaluation. Then, the body of the function is executed. On the concrete side, the execution yields a new state **mfin**. On the abstract side, the new state is obtained by running the formal model. We have to verify that the projection from the concrete state **mfin** is related to this abstract state. The proof is performed in a top-down manner. It follows the definition of the instruction, analyzing the expression step by step. The function body is split into statements and then into expressions. When evaluating an expression, one has to search for two kinds of information. The first one is how the memory state changes on the concrete side; the other is whether the results on the abstract and the concrete model are related by the projection. To this end, a library of lemmas had to be developed, identifying five categories summarized below.

1. *Evaluating a **CompCert** expression with no modification on the memory state.* Such lemmas are concerned with the expression evaluation on **CompCert C** side and in particular the C memory state change issue. Asserting that a memory state is not modified has two aspects: one is that the memory contents are not modified; the other is that the memory access permission is not changed. For example, evaluating the boolean expression `Sbit == 1` returns an unchanged memory state.

$$\text{if } G, E \vdash \text{eval\_binop}_c (Sbit == 1), M \xRightarrow{\varepsilon} v, M' \\ \text{then } M = M'.$$

In Coq syntax, the relation in premise is expressed with `eval_binop`. In this lemma and the following,  $E$  is the local environment,  $G$  is the global environment,  $M$  is the memory state,  $\varepsilon$  is the empty event (we may have here a series of events, e.g. system call, volatile load/store) and  $v$  is the result. The evaluation is performed under environments  $G$  and  $E$ . Before evaluation, we are in memory state  $M$ . With no event occurring, we get the next memory state  $M'$ . According to the definition of `eval_binop`, an internal memory state will be introduced.

$$\frac{G, E \vdash a_1, M \Rightarrow M' \quad G, E \vdash a_2, M' \Rightarrow M''}{G, E \vdash (a_1 \text{ binop } a_2), M \Rightarrow M''}$$

In the example, expression  $a_1$  is the value of `Sbit` and  $a_2$  is the constant value 1. By inverting the hypothesis of type `eval_binop`, we obtain several new hypotheses, including on the evaluation of the two subexpressions and

the introduction of an intermediate memory state  $M''$ . Evaluating them has no change on the C memory state, hence we have  $M = M'' = M'$ . In more detail, from the **CompCert** C semantics definition, we know that the evaluation of an expression will change the memory state if the evaluation contains uses of `store_value_of_type`. In **CompCert**, the basic store function on memory is represented by an inductive type `assign_loc` instead of `store_value_of_type`. As a note, since **CompCert** version supports volatile memory access, we also have to determine whether the object type is volatile before storage.

2. *Result of the evaluation of an expression with no modification on the memory.* Continuing the example above, we now discuss the result of evaluating the binary operation `Sbit == 1` both in the abstract and the concrete model. At the end of evaluation, a boolean value *true* or *false* is returned in both the concrete and the abstract models.

if `Sbit_related`  $M$  `Sbit`,  
 and  $G, E \vdash \text{eval\_rvalue\_binop}_c (\text{Sbit} == 1), M \Rightarrow v, M'$   
 then  $v = (\text{Sbit} == 1)_{coq}$

Intuitively, the projection corresponding to the parameter `sbit` in the concrete model must yield the same value as in the abstract model. Here, the expression is a so-called “simple expression” that always terminates in a deterministic way, and preserves the memory state. To evaluate the value of simple expressions, **CompCert** provides two big-step relations `eval_simple_rvalue` and `eval_simple_lvalue` for evaluating respectively their left and right values. The rules have the following shape:

$$\frac{G, E \vdash a_1, M \Rightarrow v_1 \quad G, E \vdash a_2, M \Rightarrow v_2 \quad \text{sem\_binary\_operation}(op, v_1, v_2, M) = v}{G, E \vdash (a_1 \text{ op } a_2), M \Rightarrow v}$$

In order to evaluate the binary expression  $a_1 \text{ op } a_2$ , the sub-expressions  $a_1$  and  $a_2$  are first evaluated, and their respective results  $v_1$  and  $v_2$  are used to compute the final result  $v$ .

3. *Memory state changed by storage operation or side effects.*

As mentioned before, evaluating some expressions such as `eval_assign` may modify the memory state. Lemmas are required to state that corresponding variables in the abstract and in the concrete model must evolve consistently. For example, considering an assignment on register  $Rn$ , the projection relation `register_related` is used. Expressions with side effects of modifying memory are very similar.

if `rn_related`  $M$   $rn$   
 and  $G, E \vdash \text{eval\_assign}_c (rn := rx), M \Rightarrow M', v$   
 then `rn_related`  $M'$   $rn$

4. *Internal function call.*

The simulation code is sometimes using functions from libraries. We distinguish

**internal** functions and **external** functions. An internal function is a function that belongs to a library, the code of which is part of the simulator, that we have coded ourselves, or the C code is provided by compcert C. An external function is a function for which we do not have access to the operational semantics. After an internal function is called, a new stack of blocks is typically allocated in memory. After the evaluation of the function, these blocks will be freed. Unfortunately, this may not bring the memory back to the previous state: the memory contents may stay the same, but pointers and memory organization may have changed.

if **proc\_state\_related**  $M$   $st$   
 and  $G, E \vdash \text{eval\_funcall}_c(\text{copy\_StatusRegister})_c, M \Rightarrow v, M'$   
 and  $st' = (\text{copy\_StatusRegister})_{coq} st$   
 then **proc\_state\_related**  $M'$   $st'$ .

Lemmas must be developed regarding the evaluation of internal functions, so that one can observe the returned results, compare it with the corresponding evaluation in the formal specification, and verify some conditions. For example, the lemma above is about the processor state after evaluating an internal function call `copy_StatusRegister`, which reads the value of the CPSR (Current Processor Status Register) and copies it into the SPSR (Saved Processor Status Register) when an exception occurs. The evaluation of `copy_StatusRegister` must be protected by a check on the current processor mode. If it is in authorized mode, the function `copy_StatusRegister` can be called. Otherwise, the result is “unpredictable”, which is defined by ARM architecture

It is necessary to reason on the newly returned states, which should still be related by the projection. This step is usually easy to prove, by calculation on the two representations of the processor state to verify that they match.

#### 5. External function call.

The **CompCert** C AST of an external function call contains the types of input arguments and of the returned value, and an empty body. For each external function (e.g. `memcpy()`), we have its asserted properties. mostly provided by **CompCert** C. The general expected properties of an external call are that (i) the call returns a result, which has to be related to the abstract state, (ii) the arguments must comply with the signature. (iii) after the call, no memory blocks are invalidated, (iv) the call does not increase the access permission of any valid block, and finally that the memory state can be modified only when the access permission of the call is granted. For each external call, such required properties are verified.

### 4.5 Inversion

Equipped with these lemmas we can build the proof scripts for ARM instructions. For that, we are decomposing the ARM instruction execution step by step to perform the execution of the C programs. **CompCert** C operational semantics define large and complex inductive relations. Each constructor describes the memory state transformation of an expression, statement, or function. As soon

as we want to discover the relation between memory states before and after evaluating the C code, we have to *invert* the hypotheses of operational semantics to follow the clue given by its definition, to verify the hypotheses relating concrete memory states according to the operational semantics.

An *inversion* is a kind of forward reasoning step that allows for users to extract all useful information contained in a hypothesis. It is an analysis over the given hypothesis according to its specific arguments, that removes absurd cases, introduces relevant premises in the environment and performs suitable substitutions in the whole goal. Most proof assistants provide an inversion mechanism. In the case of Coq, it is a general tactic called **inversion** [?].

Every instruction contains complex expressions, but each use of **inversion** will go one step only. If we want to find the relation between the memory states affected by these expressions, we have to invert many times. For illustration, let us consider the simple example from the ARM reference manual `CPSR = SPSR`, that assigns to register CPSR the value of SPSR (defined above). As the status register is not implemented by a single value, but a set of individual fields, the corresponding C code is a call to the function `copy_StatusRegister`, which sets the CPSR field by field with the values from SPSR. Lemma `same_cp_SR` below states that the C memory state of the simulator and the corresponding formal representation of ARM processor state evolve consistently during this assignment.

```
Lemma same_copy_SR :
  ∀ e m l b s t m' v em,
  proc_state_related m e (Ok tt (mk_semstate l b s)) →
  eval_expression (Genv.globalenv prog_adc) e m expr_cp_SR t m' v →
  ∀ l b, proc_state_related m' e
    (Ok tt (mk_semstate l b (Arm6.State.set_cpsr s
      (Arm6.State.spsr s em))))
```

In its proof, 18 consecutive inversions are needed in order to exhaust all constructors occurring in the assumptions. Unfortunately, **inversion** generates uncontrollable names which pollute proof scripts. Here, an intensive use of **inversion** makes proofs scripts unmanageable, and not robust to version changes of Coq or **CompCert**. In order to reduce the script size and get better maintainability, we studied a general solution to the inversion problem, and developed a new mechanism described in [?]. On top of it, we could program a Coq tactic able to automatically find the hypothesis to invert by matching the targeted memory states, properly manage other hypotheses, perform our inversion, clean up the goal, and repeat the above steps until all transitions between the two targeted memory states are discovered.

As a result, proofs script have become much shorter and more manageable. Considering the former example of `same_copy_SR`, the 18 calls to standard **inversion** reduce into one single step: `inv_eval_expr m m'`.

## 4.6 Instruction Proofs

Proofs of instructions rely heavily on the library of lemmas and the controllable inversion mechanism described above. Scripts size vary with the instructions complexity from less than 200 lines (e.g 170 for LDRB) to over 1000 (1204 for ADC). As a result, for each ARM instruction, we have established a theorem proving that the C code simulating an ARM instruction is equivalent to the formal specification of the ARM processor.

## 5 Conclusion

Using the approach presented in this paper, we have constructed a tool chain that makes it possible to certify that the simulation of a binary executable program on some simulation platform is compliant with the formal model of the target hardware architecture. Using CompCert-C, that has defined formal C semantics, we have formally proved, using the Coq theorem prover, the ARM v6 Instruction Set Simulator of SimSoc.

We certainly acknowledge the limits of our approach: the quality of our “verified simulation” relies on the faithfulness of our formal model of the ARM processor to the real hardware. Because the vendor companies do not provide a formal description of their hardware, one has to build them<sup>4</sup>. This issue is partly solved in this work by automatically deriving the most tedious parts of the Coq formal model from pseudo-code extracted from the vendor reference manual. If the vendors would make public formal specifications of their architectures, then our toolchain would become fully verified.

We believe this work has further impact on proofs of programs. First, we have proved here a significantly large C program. Second, because the proved program is a hardware simulator, it can be used as a tool to prove execution of target programs. For example considering a cryptographic algorithm implemented for the ARM architecture and compiled with CompCert-C, it could then be proved that the execution of that program provides the exact encryption required, and nothing else. Therefore, the tool presented is an enabler for the proofs of other programs, which offers a direction for future research.

Another consequence of this work is that, supposing one could compile the C instructions to silicon using a silicon compiler, and that compiler would also be certified, ala **CompCert**, it would then make it possible to prove real hardware...

---

<sup>4</sup> Note that this problem is the same as for the work done by Cambridge University.